

simulation_MG1_800G

2024-11-29

```
source("func/igraph_functions.R")
source("func/general_function_base.R")
source("func/simmer_function_base.R")
```

#Input data

The **average packet size** is taken from: Amsterdam Internet Exchange Ethernet Frame Size Distribution, statistics available online at <https://stats.ams-ix.net/sflow/size.html>, accessed on July 2023

Packet sizes in Bytes: 64-127, 128-511, 512 - 1023, 1024 - 1513, 1514, more than 1515.

With their probabilities: 0.332, 0.054, 0.033, 0.037, 0.346, 0.146, 0.052

```
PS_size=c((64+127)/2,(128+255)/2,(256+511)/2, (512+1023)/2, (1024+1513)/2, 1514, (1515+9100)/2)
PS_weights=c(33.2/100, 5.4/100, 3.3/100, 3.7/100, 34.6/100, 14.6/100, 5.2/100)
N = sum(PS_size*PS_weights)
N
```

```
## [1] 1019.035
```

```
var_N <- sum(PS_size^2*PS_weights) - N^2
Cs2 <- var_N/(N^2)

CapacityGbps = 800
#Calculation of delays for different capacities
Load_local = 0.3
Load_regional = 0.5
Load_national = 0.4
```

Input excel file with topology information: links, nodes, traffic

Prefix base - the common name for the nodes in the file.

```
#topology_choice <- readline(prompt = "Enter 'Tokyo' or 'Milano' to choose the respective topology: ")

# # Define the file paths based on user input
# file_name_v2 <- "input_files/Metro_topology_full_Tokyo.xlsx"
# prefix_base <- "Tokyo_"
# topology_name = "Tokyo"
# print(file_name_v2)

file_name_v2 <- "input_files/Metro_topology_full_Milano.xlsx"
prefix_base <- "Node"
topology_name = "Milano"
print(file_name_v2)
```

```
## [1] "input_files/Metro_topology_full_Milano.xlsx"
```

```
#  
# file_name_v2 <- "input_files/Metro_topology_MAN157.xlsx"  
# prefix_base <- ""  
# topology_name = "MAN157"  
# print(file_name_v2)
```

```
nodes_info <- read_excel(file_name_v2, sheet = 1)  
links_info <- read_excel(file_name_v2, sheet = 2)  
traffic_file <- read_excel(file_name_v2, sheet = 3)
```

```
print(links_info)
```

```
## # A tibble: 202 x 4  
##   sourceID destinationID distanceKm capacityGbps  
##   <chr>      <chr>          <dbl>      <dbl>  
## 1 Node1     Node2              1          0  
## 2 Node1     Node3              1          0  
## 3 Node1     Node4              1          0  
## 4 Node1     Node8              1      3824.  
## 5 Node2     Node1              1          0  
## 6 Node2     Node5              1     3998.  
## 7 Node2     Node6              1          0  
## 8 Node2     Node7              1          0  
## 9 Node3     Node1              1          0  
## 10 Node3    Node4              0.8         0  
## # i 192 more rows
```

```
print(nodes_info)
```

```
## # A tibble: 52 x 11  
##   node_name node_code 'Node Type' 'Central office type' Reference Regional C-1  
##   <chr>      <chr>      <chr>      <chr>          <chr>  
## 1 Node1     HL2        Metro Core Regional CO      Node1  
## 2 Node2     HL2        Metro Core Regional CO      Node2  
## 3 Node3     HL2        Metro Core Regional CO      Node3  
## 4 Node4     HL2        Metro Core Regional CO      Node4  
## 5 Node5     HL2        Metro Core Regional CO      Node5  
## 6 Node6     HL2        Metro Core Regional CO      Node6  
## 7 Node7     HL2        Metro Core Regional CO      Node7  
## 8 Node8     HL2        Metro Core Regional CO      Node8  
## 9 Node9     HL3        Metro Core ~ National CO      Node8  
## 10 Node10    HL3        Metro Core ~ National CO      Node8  
## # i 42 more rows  
## # i abbreviated name: 1: 'Reference Regional CO'  
## # i 6 more variables: 'Reference National CO' <chr>, Households <dbl>,  
## #   'Macro cells sites' <dbl>, 'Small cell sites' <dbl>,  
## #   'Twin Regional CO' <chr>, 'Twin National CO' <chr>
```

```
print(traffic_file)
```

```
## # A tibble: 104 x 4
##   sourceID destinationID trafficGbps service
##   <chr>      <chr>          <dbl> <chr>
## 1 Node1      Node9             1945. CWB
## 2 Node2      Node16            1671. CWB
## 3 Node3      Node11            1498. CWB
## 4 Node4      Node13            3015. CWB
## 5 Node5      Node16            1805. CWB
## 6 Node6      Node19            3565. CWB
## 7 Node7      Node9             2618. CWB
## 8 Node8      Node9             1916. CWB
## 9 Node9      Node9             3798. CWB
## 10 Node10    Node10            3911. CWB
## # i 94 more rows
```

Definition of national and regional Cental Offices (COs).

```
national_nodes <- c()
regional_nodes <- c()

for (i in seq_along(nodes_info$node_code)) {
  if (nodes_info$node_code[i] == "HL2") {
    national_nodes <- c(national_nodes, i)
  }
  if (nodes_info$node_code[i] == "HL3") {
    regional_nodes <- c(regional_nodes, i)
  }
}

cat("National nodes:", national_nodes, "\n")
```

```
## National nodes: 1 2 3 4 5 6 7 8 21 22 23 24
```

```
cat("Regional nodes:", regional_nodes, "\n")
```

```
## Regional nodes: 9 10 11 12 13 14 15 16 17 18 19 20 40 41 42 43 44
```

Functions

Vysochanskij–Petunin’s bound calculation function

Inputs: - **delay_hops**: Vector of delay at each hop - **a**: Upper bound percentile from 0 to 1

Output: - Vysochanskij–Petunin’s upper bound

```
func_bounds_VP <- function(delay_hops, a)
{
  mu = sum(delay_hops)
```

```

sigma = sqrt(sum(delay_hops^2))
k <- sqrt(4/9/(1-a))
Prop_VP <- 1 - 4/9/(k^2)
upper_bound_VP <- k * sigma + mu
return(upper_bound_VP)
}

```

##Simmer simulation

Igraph calculations

Building the graph:

```
g <- graph_from_data_frame(links_info, directed = TRUE, vertices = nodes_info)
```

Calculations of the capacity in p/s

```

E(g)$Distance <- E(g)$distanceKm
E(g)$Definition <- paste0(as_edgelist(g)[,1], ">", as_edgelist(g)[,2])
E(g)$Capacity <- E(g)$capacityGbps*10^9/(8*N)

```


Plot graph

```

#plot graph
V(g)$color <- "tomato"
V(g)$color[national_nodes] <- "yellow"

deg <- degree(g, mode="all")
V(g)$size <- deg*1.5
l <- layout_nicely(g)
plot(g, edge.arrow.size=.3, vertex.label = V(g)$name, edge.curved=.5, layout=l)

```



```
#Igraph calculations ## Total traffic calculatons ## Calculations of the load, average number of packets,
queuing and propagation delays: Load for Local COs 0.3; Regional COs 0.5; Rational COs 0.4
```

```
c(g_c10, data_av, data_99) := simulation_igraph(nodes_info, links_info, traffic_file, Capacity = 10, ca
```

```
#Simmer simulation Including queuing, transmission, propagation delay Output table with links information
```

```
Simulation of all traffic flows, comparision of experimental with theretical
```

```
#Results
```

```
links_info_df
```

```
## # A tibble: 202 x 5
##   sourceID destinationID distanceKm capacityGbps prop_delay_s
##   <chr>      <chr>          <dbl>      <dbl>      <dbl>
## 1 Node1      Node2              1          0      0.000005
## 2 Node1      Node3              1          0      0.000005
## 3 Node1      Node4              1          0      0.000005
## 4 Node1      Node8              1      3824.      0.000005
## 5 Node2      Node1              1          0      0.000005
## 6 Node2      Node5              1      3998.      0.000005
## 7 Node2      Node6              1          0      0.000005
## 8 Node2      Node7              1          0      0.000005
## 9 Node3      Node1              1          0      0.000005
## 10 Node3     Node4              0.8          0      0.000004
## # i 192 more rows
```

```
traffic_file_v2
```

```
## # A tibble: 72 x 11
##   sourceID destinationID trafficGbps service traffic_ps latencyPropTransQueui~1
##   <chr>      <chr>          <dbl> <chr>          <dbl>          <dbl>
## 1 Node1      Node9             240 CWB          29439617.      0.00000744
## 2 Node2      Node16            240 CWB          29439617.      0.00000744
## 3 Node3      Node11            240 CWB          29439617.      0.00000337
## 4 Node4      Node13            240 CWB          29439617.      0.00000337
## 5 Node5      Node16            240 CWB          29439617.      0.00000337
## 6 Node6      Node19            240 CWB          29439617.      0.00000337
## 7 Node7      Node9             240 CWB          29439617.      0.00000694
## 8 Node8      Node9             240 CWB          29439617.      0.00000337
## 9 Node20     Node19            240 CWB          29439617.      0.00000337
## 10 Node21    Node40            240 CWB          29439617.      0.00000337
## # i 62 more rows
## # i abbreviated name: 1: latencyPropTransQueuing_theor_s_mg1
## # i 5 more variables: latencyPercentile99_theor_s_mg1 <dbl>, Delay_sim_s <dbl>,
## #   Delay_th_s <dbl>, D99_sim_s <dbl>, VPbound_99th_s <dbl>
```

```
#Calculation of delays for different capacities
```

```
capacities = c(400, 800, 1200, 1600, 3200)
c(g_c400_dg_l05, data_av_c400_dg_l05, data_99_c400_dg_l05, trafficGbps_c400_dg_l05) :=simulation_igraph
c(g_c800_dg_l05, data_av_c800_dg_l05, data_99_c800_dg_l05, trafficGbps_c800_dg_l05) :=simulation_igraph
c(g_c12_dg_l05, data_av_c12_dg_l05, data_99_c12_dg_l05, trafficGbps_c12_dg_l05) :=simulation_igraph(nod
c(g_c16_dg_l05, data_av_c16_dg_l05, data_99_c16_dg_l05, trafficGbps_c16_dg_l05) :=simulation_igraph(nod

c(g_c32_dg_l05, data_av_c32_dg_l05, data_99_c32_dg_l05, trafficGbps_c32_dg_l05) :=simulation_igraph(nod
```

```
text_size = 12
```

```
# Boxplot for trafficGbps (bit rate in Gbps)
df_plot_traffic <- data.frame(capacity = c(
  rep("400G", length(trafficGbps_c400_dg_l05)),
  rep("800G", length(trafficGbps_c800_dg_l05)),
  rep("1200G", length(trafficGbps_c12_dg_l05)),
  rep("1600G", length(trafficGbps_c16_dg_l05)),
  rep("3200G", length(trafficGbps_c32_dg_l05))),
  value = c(trafficGbps_c400_dg_l05,
    trafficGbps_c800_dg_l05,
    trafficGbps_c12_dg_l05,
    trafficGbps_c16_dg_l05,
    trafficGbps_c32_dg_l05),
  type_of_calc = "bit rate",
  city = c(rep(topology_name,length(trafficGbps_c32_dg_l05)*10)))

ggplot(df_plot_traffic, aes(x = fct_reorder(capacity, value, .desc = TRUE), y = value, fill = type_of_c
  geom_boxplot() +
  theme_bw() +
  theme(
    legend.title = element_blank(),
    legend.position = c(0.8, 0.99),
```

```

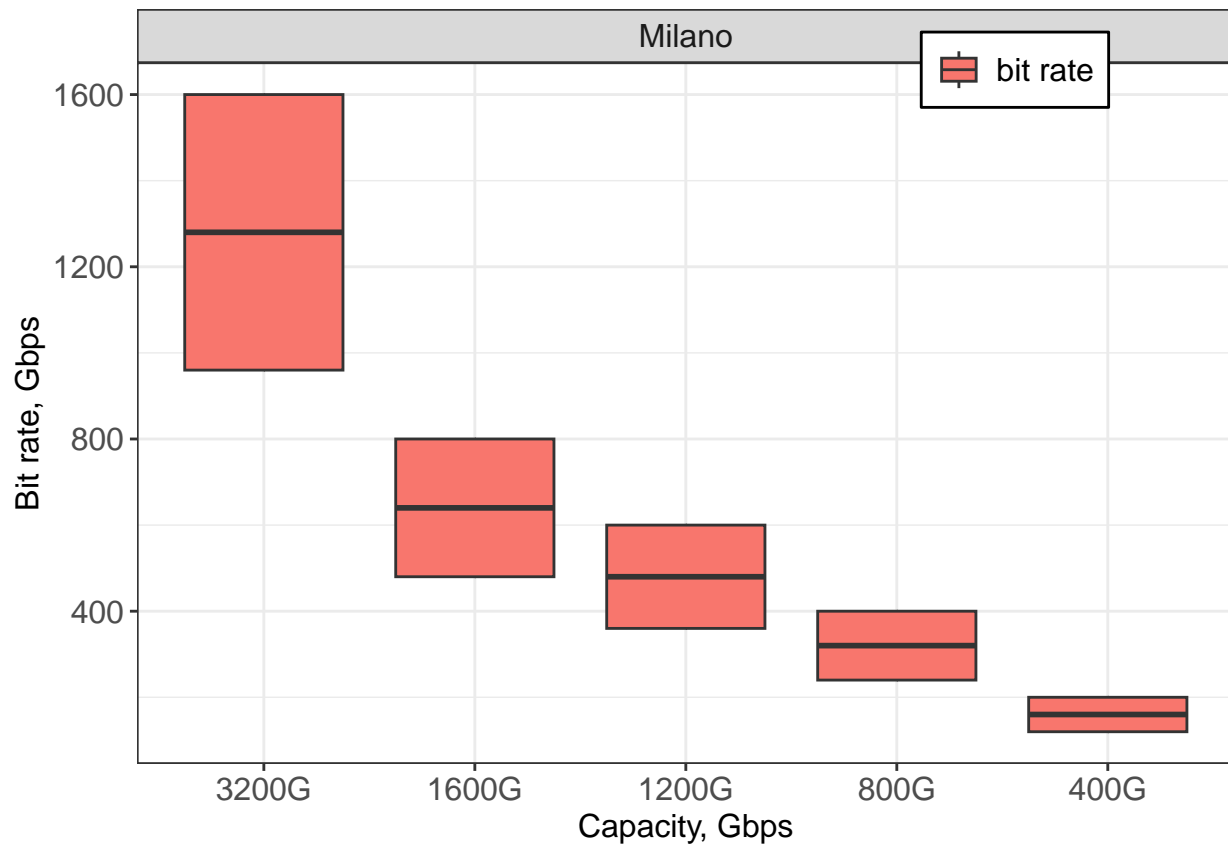
axis.title.y = element_text(size = text_size),
axis.title.x = element_text(size = text_size),
legend.text = element_text(size = text_size),
legend.background = element_rect(colour = "black"),
axis.text = element_text(size = text_size),
strip.text = element_text(size = text_size)
) +
facet_grid(. ~ city ) +
ylab("Bit rate, Gbps") +
xlab("Capacity, Gbps")

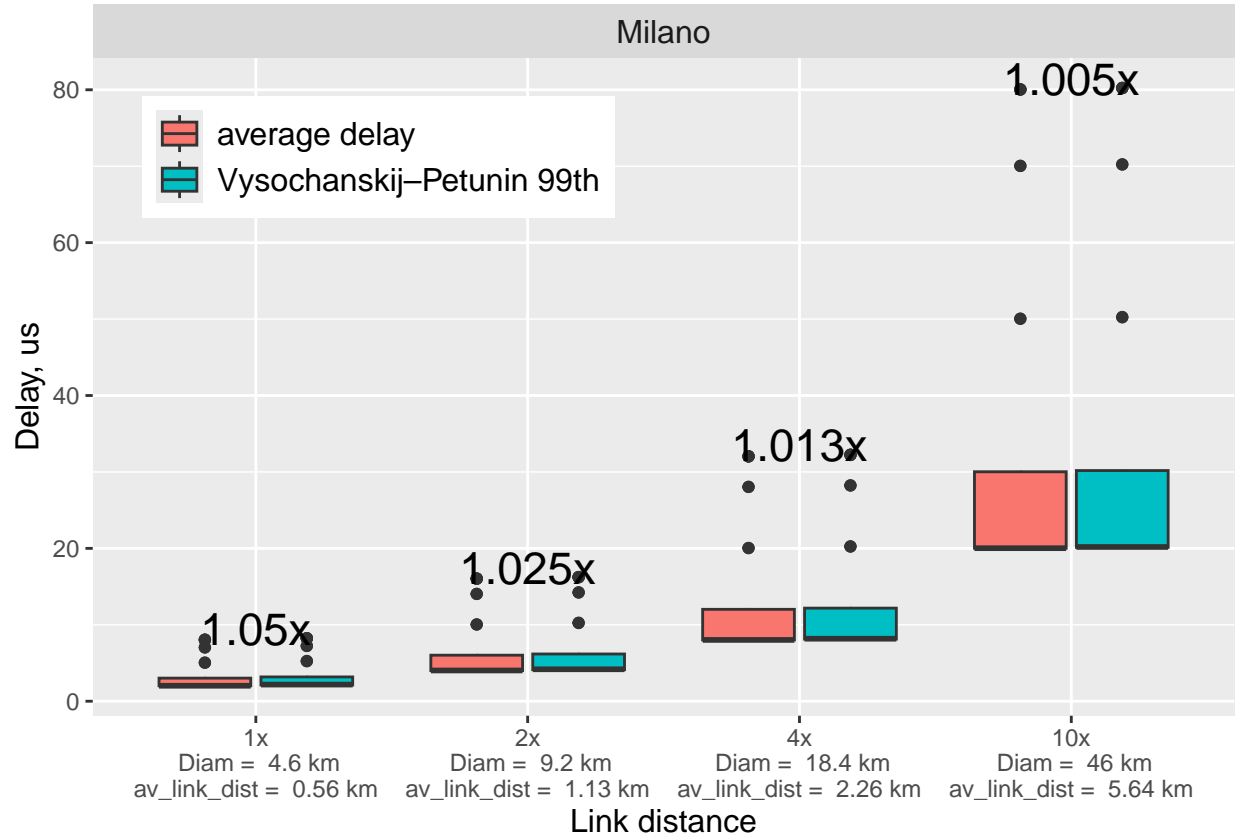
```

```

## Warning: A numeric 'legend.position' argument in 'theme()' was deprecated in ggplot2
## 3.5.0.
## i Please use the 'legend.position.inside' argument of 'theme()' instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.

```





```
prop_queue_tr_delay_e2e <- function(g, traffic_file, Cs2){

  E(g)$Queue_Delay <- ifelse(E(g)$Traffic == 0, 0, E(g)$Load/(1-E(g)$Load)/E(g)$Traffic*(Cs2 + 1)/2) #M
  latencyPropTransQueuing_theor_s <- rep(0, nrow(traffic_file))
  latencyPercentile99_theor_s <- rep(0, nrow(traffic_file))

  latencyProp_s <- rep(0, nrow(traffic_file))
  latencyPercentile99_theor_TransQueuing_s <- rep(0, nrow(traffic_file))
  processing_FEC_s <- rep(0, nrow(traffic_file))

  for (line in 1:nrow(traffic_file))
  {
    if(traffic_file$traffic_ps[line] != 0){
      vertex_source <- which(V(g)$name==traffic_file$sourceID[line])
      vertex_destination<- which(V(g)$name==traffic_file$destinationID[line])
      path <- shortest_paths(g, vertex_source, vertex_destination,
                            weights = NULL,
                            output = "both",
                            algorithm = c("automatic"))
      latencyPropTransQueuing_theor_s[line] = sum(E(g)[path[["epath"]][[1]]]$Queue_Delay) + sum(E(g)[pa
      latencyPercentile99_theor_s[line] <- func_bounds_VP(E(g)[path[["epath"]][[1]]]$Queue_Delay, 0.99)

      latencyProp_s[line] = sum(E(g)[path[["epath"]][[1]]]$Prop_Delay)
      latencyPercentile99_theor_TransQueuing_s[line] <- func_bounds_VP(E(g)[path[["epath"]][[1]]]$Queue
      processing_FEC_s[line] <- 4*1e-6*length(path[["epath"]][[1]])
    }
  }
}
```

```

    else {
      latencyPropTransQueuing_theor_s[line] <- 0
      latencyPercentile99_theor_s[line] <- 0
      latencyProp_s[line] <- 0
      latencyPercentile99_theor_TransQueuing_s[line] <- 0
      processing_FEC_s[line] <- 0
    }
  }
}

return(list(g, latencyPropTransQueuing_theor_s, latencyPercentile99_theor_s, latencyProp_s, latencyPe
}

simulation_igraph <- function(nodes_info, links_info, traffic_file, Capacity = 10, calc_dist = FALSE, c

  traffic_file$traffic_ps <- traffic_file$trafficGbps*10^9/(8*N)
  g = graph_from_data_frame(links_info, directed = TRUE)
  E(g)$capacityGbps <- Capacity
  for (NCO in national_nodes){
    a <- (filter(traffic_file, traffic_file$destinationID == V(g)$name[NCO]))
    size <- length(a$destinationID)
    E(g)$capacityGbps[incident(g, NCO, mode = c("in"))] <- Capacity*size
  }
  E(g)$Capacity <- E(g)$capacityGbps*1e9/(8*N) #p/s
  #####distance put original distance(calc_dist = 1) or custom distance (calc_dist = 1) with par
  if( calc_dist == TRUE ) {
    E(g)$Distance = E(g)$distanceKm * distance
  }
  else {
    E(g)$Distance <- E(g)$distanceKm
  }
  ####Load of the links
  #for Local CO
  E(g)$Load <- Load_local
  #national COs
  for (NCO in national_nodes){
    E(g)$Load[incident(g, NCO, mode = c("in"))] <- Load_national
  }
  #regional COs
  for (RCO in regional_nodes){
    E(g)$Load[incident(g, RCO, mode = c("in"))] <- Load_regional
  }

  E(g)$Traffic <- E(g)$Capacity * E(g)$Load
  E(g)$Ni = E(g)$Load/(1-E(g)$Load) # average number of packets in each system
  E(g)$Prop_Delay <- 5*10^(-6)*E(g)$Distance
  c(g, traffic_file$latencyPropTransQueuing_theor_s_mg1, traffic_file$latencyPercentile99_theor_s_mg1,

  data_prop <- latencyProp_s_mg1
  data_99_TransQueuing <- latencyPercentile99_theor_TransQueuing_s_mg1

```

```

assign("traffic_file", traffic_file, .GlobalEnv)

links_info$prop_delay_s <- 5*10(-6)*links_info$distanceKm
assign("links_info_df", links_info, .GlobalEnv)

#removing the local traffic
data_av <- data_av[data_av != 0]
data_99 <- data_99[data_99 != 0]
trafficGbps <- E(g)$Capacity * E(g)$Load[E(g)$Capacity * E(g)$Load != 0]/1e9*(8*N) #Gbps

return(list(g, data_prop, data_99_TransQueueing, trafficGbps, processing_FEC_s))

}

```

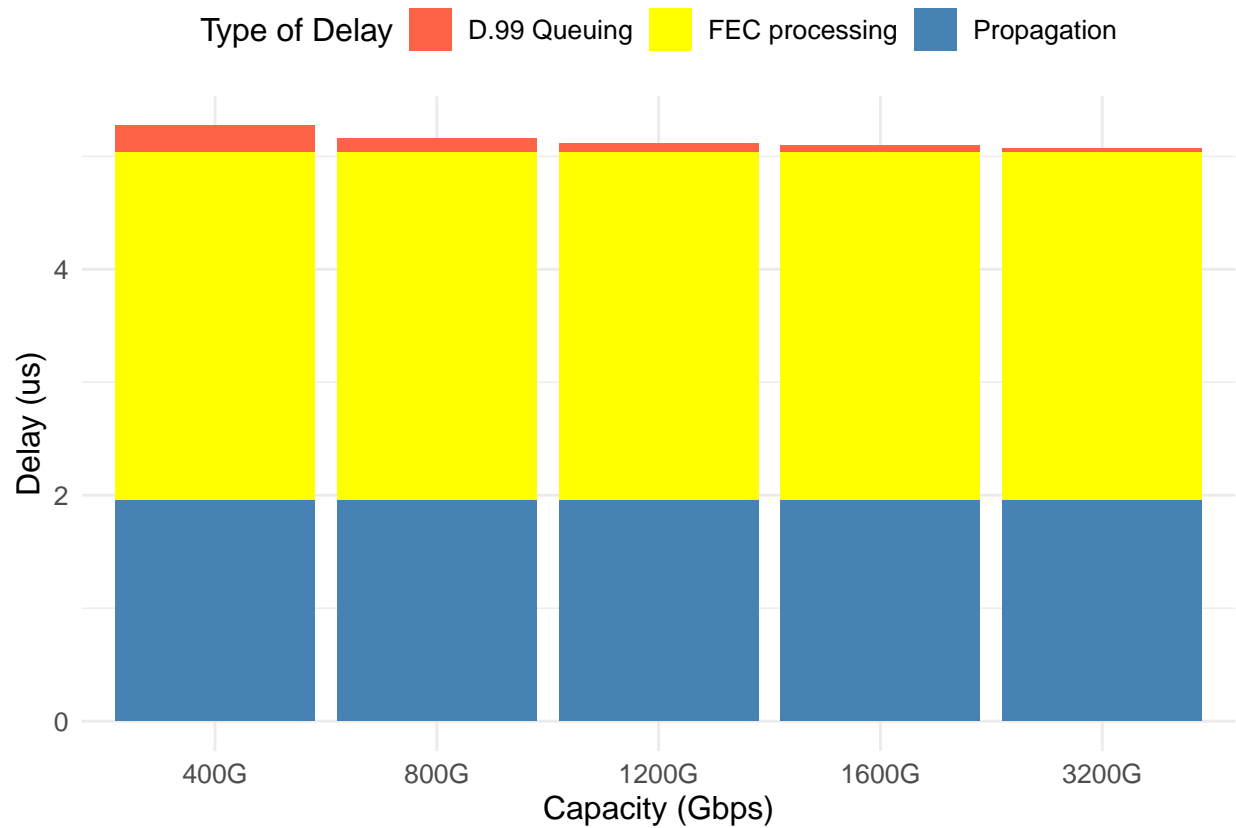
```

capacities = c(400, 800, 1200, 1600, 3200)
c(g_c400_dg_l05, data_prop_c400_dg_l05, data_99_TransQueueing_c400_dg_l05, trafficGbps_c400_dg_l05, processing_FEC_s_c400_dg_l05)
c(g_c800_dg_l05, data_prop_c800_dg_l05, data_99_TransQueueing_c800_dg_l05, trafficGbps_c800_dg_l05, processing_FEC_s_c800_dg_l05)
c(g_c12_dg_l05, data_prop_c12_dg_l05, data_99_TransQueueing_c12_dg_l05, trafficGbps_c12_dg_l05, processing_FEC_s_c12_dg_l05)
c(g_c16_dg_l05, data_prop_c16_dg_l05, data_99_TransQueueing_c16_dg_l05, trafficGbps_c16_dg_l05, processing_FEC_s_c16_dg_l05)
c(g_c32_dg_l05, data_prop_c32_dg_l05, data_99_TransQueueing_c32_dg_l05, trafficGbps_c32_dg_l05, processing_FEC_s_c32_dg_l05)

# Create a data frame with mean propagation and queueing delays for each capacity
df_plot_delay <- data.frame(
  capacity = factor(c("400G", "800G", "1200G", "1600G", "3200G"),
    levels = c("400G", "800G", "1200G", "1600G", "3200G")),
  delay_type = rep(c("Propagation", "D.99 Queueing", "FEC processing"), each = 5),
  delay_value = c(
    mean(data_prop_c400_dg_l05), mean(data_prop_c800_dg_l05), mean(data_prop_c12_dg_l05), mean(data_prop_c16_dg_l05), mean(data_prop_c32_dg_l05),
    mean(data_99_TransQueueing_c400_dg_l05), mean(data_99_TransQueueing_c800_dg_l05), mean(data_99_TransQueueing_c12_dg_l05), mean(data_99_TransQueueing_c16_dg_l05), mean(data_99_TransQueueing_c32_dg_l05),
    mean(processing_FEC_s_c400_dg_l05), mean(processing_FEC_s_c800_dg_l05), mean(processing_FEC_s_c12_dg_l05), mean(processing_FEC_s_c16_dg_l05), mean(processing_FEC_s_c32_dg_l05)
  )
)

# Plot the barplot
ggplot(df_plot_delay, aes(x = capacity, y = delay_value * 1e6, fill = delay_type)) +
  geom_bar(stat = "identity", position = "stack") +
  theme_minimal() +
  labs(
    x = "Capacity (Gbps)",
    y = "Delay (us)",
    fill = "Type of Delay"
  ) +
  scale_fill_manual(values = c("Propagation" = "steelblue", "D.99 Queueing" = "tomato", "FEC processing" = "firebrick")) +
  theme(
    text = element_text(size = 12),
    legend.position = "top"
  )

```



```
df_Milano = data.frame(prop = c(mean(data_prop_c400_dg_l05), mean(data_prop_c800_dg_l05), mean(data_prop_c1200_dg_l05), mean(data_prop_c1600_dg_l05), mean(data_prop_c3200_dg_l05)),
                        queuing = c( mean(data_99_TransQueuing_c400_dg_l05), mean(data_99_TransQueuing_c800_dg_l05), mean(data_99_TransQueuing_c1200_dg_l05), mean(data_99_TransQueuing_c1600_dg_l05), mean(data_99_TransQueuing_c3200_dg_l05)))
```

```
#Creating the report file
```

```
update_file(links_info_df, traffic_file_v2, file_name = paste0(topology_name, "_topology_result.xlsx"),
```